

Weight Prediction with the help of Gender and Height

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
import warnings
warnings.simplefilter("ignore")
```

```
In [2]: dataset = pd.read_csv("weight-height.csv")
```

```
In [3]: dataset.head()
```

```
Out[3]:
```

	Gender	Height	Weight
0	Male	73.847017	241.893563
1	Male	68.781904	162.310473
2	Male	74.110105	212.740856
3	Male	71.730978	220.042470
4	Male	69.881796	206.349801

```
In [4]: dataset.shape
```

```
Out[4]: (10000, 3)
```

```
In [5]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Gender   10000 non-null    object
1   Height   10000 non-null    float64
2   Weight   10000 non-null    float64
dtypes: float64(2), object(1)
memory usage: 234.5+ KB
```

```
In [6]: dataset.describe()
```

```
Out[6]:
```

	Height	Weight
count	10000.000000	10000.000000
mean	66.367560	161.440357

	Height	Weight
std	3.847528	32.108439
min	54.263133	64.700127
25%	63.505620	135.818051
50%	66.318070	161.212928
75%	69.174262	187.169525
max	78.998742	269.989699

In [7]:

dataset.count()

Out[7]: Gender 10000
Height 10000
Weight 10000
dtype: int64

In [8]:

dataset.isnull().sum()

Out[8]: Gender 0
Height 0
Weight 0
dtype: int64

In [9]:

dataset["Gender"] = LabelEncoder().fit_transform(dataset["Gender"])

In [10]:

dataset.head()


Out[10]:

	Gender	Height	Weight
0	1	73.847017	241.893563
1	1	68.781904	162.310473
2	1	74.110105	212.740856
3	1	71.730978	220.042470
4	1	69.881796	206.349801

In [11]:

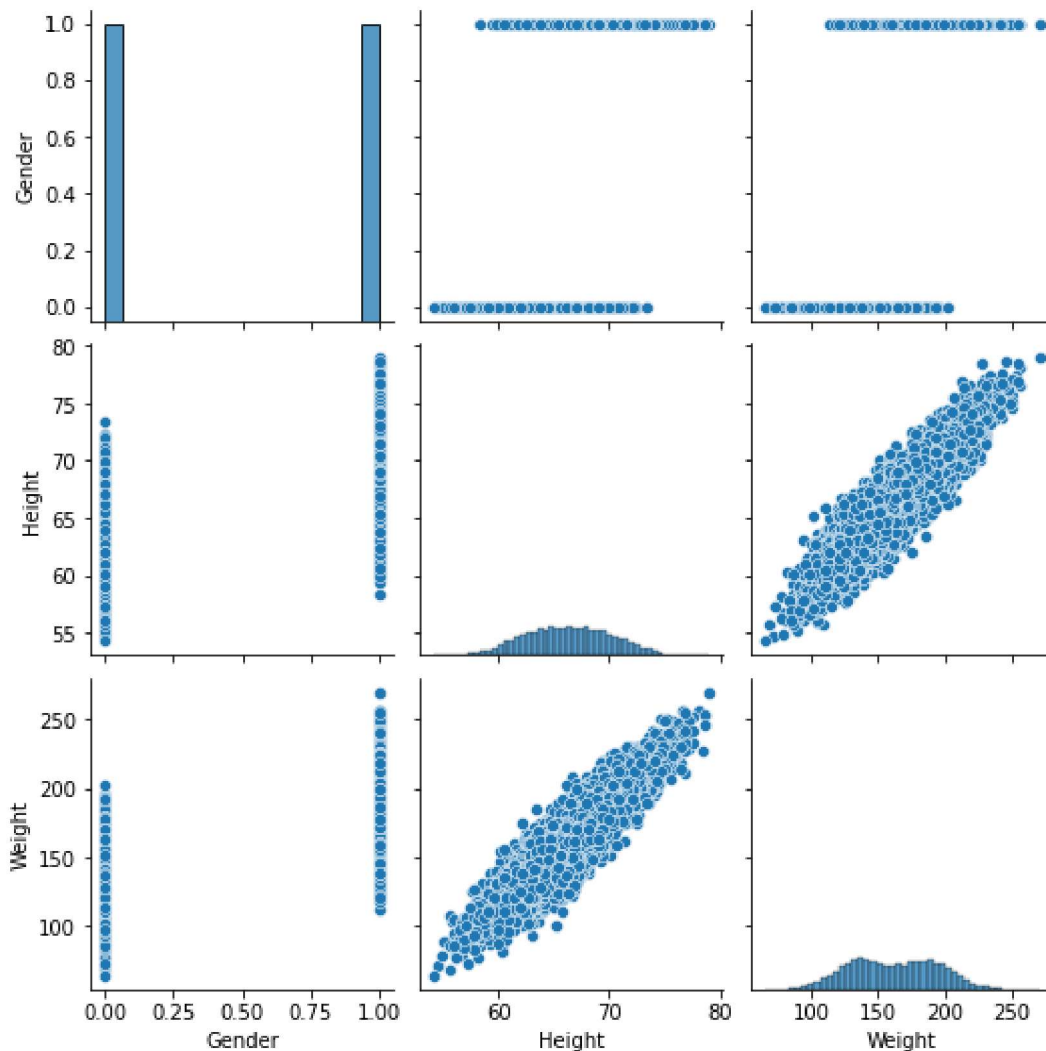
sns.countplot(x="Gender", data=dataset)

Out[11]: <AxesSubplot:xlabel='Gender', ylabel='count'>



```
In [12]: sns.pairplot(dataset)
```

```
Out[12]: <seaborn.axisgrid.PairGrid at 0x20e19430fa0>
```



```
In [13]: dataset.columns
```

```
Out[13]: Index(['Gender', 'Height', 'Weight'], dtype='object')
```

```
In [14]: X = dataset.drop(columns="Weight",axis=1)
          Y = dataset["Weight"]
```

In [15]:

X.head()

Out[15]:

	Gender	Height
0	1	73.847017
1	1	68.781904
2	1	74.110105
3	1	71.730978
4	1	69.881796

In [16]:

X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.25,random_state=0)

In [17]:

X_train.head()

Out[17]:

	Gender	Height
2967	1	68.058837
700	1	69.760095
3481	1	71.702360
1621	1	71.096113
800	1	72.215035

In [18]:

X_test.head()

Out[18]:

	Gender	Height
9394	0	64.723877
898	1	67.272171
2398	1	74.520972
5906	0	65.880014
2343	1	64.846301

In [19]:

Y_train.head()

Out[19]:

2967	187.779075
700	187.812062
3481	214.787698
1621	210.821194
800	204.937760

Name: Weight, dtype: float64

In [20]:

Y_test.head()

Out[20]:

9394	138.085796
898	187.363366
2398	216.533191

```
5906      131.761443
2343      157.718438
Name: Weight, dtype: float64
```

```
In [21]: X.shape,X_test.shape,X_train.shape
```

```
Out[21]: ((10000, 2), (2500, 2), (7500, 2))
```

```
In [22]: model = LinearRegression()
```

```
In [23]: model.fit(X_train,Y_train)
```

```
Out[23]: ▾ LinearRegression
LinearRegression()
```

```
In [24]: prediction = model.predict(X_test)
```

```
In [25]: score = r2_score(prediction,Y_test)
```

```
In [26]: score
```

```
Out[26]: 0.8905030905575755
```

```
In [27]: new_data = [[0,74]]
weight = model.predict(new_data)
weight
print(weight*0.45)
```

```
[88.77203159]
```

```
In [ ]:
```